

Microsoft  
Technology  
Centers

## Data Literacy Workshop Series

Dave Wentzel

[linkedin.com/in/dwentzel](https://www.linkedin.com/in/dwentzel)

Philadelphia MTC

# Dave Wentzel



## **Decision Architect**

Microsoft Technology Center

[davew@microsoft.com](mailto:davew@microsoft.com)

[linkedin.com/in/dwentzel](https://linkedin.com/in/dwentzel)

Deloitte and Touche 7 years

Siemens Medical Solutions 5 years

Independent Consultant 4 years

CTO for data consultancy 2 years

Microsoft 4 years



# Our People

The Architects of the MTC deliver immersive industry experiences and deep technical engagement focused on business outcomes. They help you understand the art of the possible and make it real by creating solutions to achieve business outcomes.

## Architects



**Todd Furst**

*Chief Technical Architect*

Cross-workload specializing in *Azure-Infrastructure*

Industry: Retail, Healthcare




**Rich Ross**

*Principal Technical Architect*

Cross-workload specializing in *Azure-App Dev*

Industry: Healthcare, Manufacturing



**Dave Wentzel**

*Principal Technical Architect*

Data and Analytics




**Michael Mukalian**

*Principal Technical Architect*

Cross-workload specializing in *Modern Workplace*

Industry: Financial Services, Healthcare



**James Stento**

*Principal Technical Architect*

Cross-workload specializing in *Biz Apps*

Industry: Media & Communications, Manufacturing

Understanding technology is less important than understanding data

- Pick the right tool for the user and use case

Self-service analytics initiatives are “underwhelming”

Your users’ level of data literacy (the ability to find, work with, analyze, and “discuss” data is critical to building a self-service, *insights-driven* culture

These sessions (add’l 6-30 planned for 2022) are both tactical and strategic

- CoE → Prescriptive Analytics

*It is my ambition to help you better integrate business analytics into the decision-making process, and brandish it for competitive advantage.*

# Our Process

Delivering the Right Experiences for our Customers



## Offerings



### Design Thinking workshop

Explores the impact of digital transformation and innovation to help customers with vision-setting, strategy, roadmaps and organizational alignment.




### Strategy briefing

A strategic business and technical discussion to gain understanding of customer goals and challenges. Align Microsoft capabilities and solutions.



### Architecture design session

Synthesizes the business and technical requirements for a solution including an initial scope and a high-level architecture to drive next steps.



### Hackathon

A hands-on, intensely collaborative and inclusive sprint to determine the applicability of specific technologies against a set of business use cases.



### Hands-on lab

A hands-on, immersive education experience to provide the skills and familiarity of a technology to enable solution development and adoption.



### Rapid prototype

A tailored hands-on experience to demonstrate the key technical capabilities of a solution and address any challenges to accelerate decision making.

<b>Customer Behavioral Analytics</b>	Monday Feb 28 1-3pm EST	Without customers, your company is out-of-business. The most successful companies are using customer behavior data to make key business decisions. Many companies struggle with things like customer segmentation, loyalty, understanding Customer Lifetime Value, and “optimizing” churn. I’ll bet your company has the data to tell which customers have churned and which might, but most struggle with “ok, now what do I do?” As data professionals, we are uncomfortable making opinionated recommendations on what we should do next based on what the data is telling us. In this session I’ll show you how to use data and analytics processes to understand customer analytics issues and how to help your business leaders interpret their data to answer the question: “What should we do next?”.
<b>Marketing Analytics Using Demand Signal Repositories</b>	Monday March 7 1-3pm EST	Most marketing analytics is still done in Excel. The reason is simple: Your marketing team wants to do complex analytics, but it’s too cumbersome to do in your data warehouse. There’s a better way. Data analysts need to be able to guide their marketing teams using Prescriptive Analytics techniques to show them “what do we do next”. In this session we’ll look at how to enable marketing teams to do “self-service analytics” (with the data team’s guidance) by looking at real use cases like calculating Facebook campaign ROIs, performing RFM analytics, and understanding CAC (customer acquisition costs) in relation to churn, customer lifetime value, and product development. We want to show you how to build a Demand Signal Repository that goes viral inside your organization. This will be a hands-on session where we can create opinions using real data.

<b>ETL is why Analytics Projects Fail: Here's a Better Way</b>	Monday March 14 1-3pm EST	Traditional data projects spend a lot of time and resources copying data around the data ecosystem. This isn't fun for your data team, and it certainly doesn't add much value IF there is an alternative. In this session we'll show you some approaches to quickly acquire data sets, determine if they provide lift, and transform that data to add business value, quickly. You can leverage these techniques today to shorten time-to-value for any analytics project.
<b>Latent Data Analytics</b>	Monday March 21 1-3pm EST	It's likely that the most interesting data in your company is not being leveraged in your analytics. The fact is: the most valuable data your customer has is likely locked in unstructured data like Word docs and pdfs. In this session we'll show you how to crack that data and structure it to gain insights. We'll give you lots of use cases and ideas to spark your creativity.
<b>Build the Corporate Decision Factory</b>	Monday March 28 1-3pm EST	You've decided to be a more data-driven company. You are building data lakes, knowledge graphs, and data catalogs. You want to create better decision-making capabilities and focus less on the HOWs and more on the WHATs and WHYs but your team is spending a lot of time focusing on implementation bottlenecks that are derailing your transformation. At the Microsoft Technology Center we have concrete, repeatable processes that we've learned from helping our customers on their analytics journey. Simple questions like "how do I structure my analytics sandbox?" or "how can I store PII in a secure and compliant manner?" take months for many teams to implement. We want to share some patterns that you can use to shorten time-to-value and put your focus back on business outcomes.

# The Problem



“Our people aren’t ready (or capable) of doing advanced analytics”

“We are being asked to do advanced analytics by our business, but we don’t know how”

“ETL is causing our analytics project to run late”

“Our data is a mess; we need to clean it up before we do anything else.”

## What we hear from IT ...

“Our last EDW project failed”

“We don’t have data scientists that can really mine our data.”

“We’d like to integrate social media and 3<sup>rd</sup> party data, but don’t know how”

“We can’t align our data strategy with our business strategy”

“We have a self-service analytics initiative, but it isn’t working”

# What we hear from business users ...

“I want to control my own data”

“The data warehouse data is nightly; I need real-time data”

“We have lots of data, but don’t know what to do with it”

“We want to do analytics in *our* system.”


“We don’t know how to democratize data and make it an enterprise-wide capability”

“I just want to get my job done”

“The data warehouse doesn’t answer my questions”

“Even with the data warehouse and reports, I do analysis in spreadmarts”


# Analytics Maturity Models



Is this how you do  
it?

# Legacy Thinking

**The Philosophy:** Model data » Transform data » Load data » *Understand* data





---

# How to draw an Owl.

---

*"A fun and creative guide for beginners"*

---




Fig 1. Draw two circles



Fig 2. Draw the rest of the damn Owl

---

# Data Projects have a high fail rate

---


Too much time is spent in:

- Requirements gathering
- Data modeling
- ETL

Users only see the fruits of the endeavor after the reports are created

In 2014, the Project Management Institute (PMI) released its Pulse of the Profession report. PMI found that “37 percent of all organizations reported inaccurate requirements as the primary reason for project failure.”


<https://www.pmi.org/learning/library/poor-requirements-management-source-failed-projects-9341>



# A Different Approach

# Data Sandboxing/EDA

- A robust and well-proven methodology.
- Data science-like.
- Iterative.
- Stresses up-front understanding of data.
- Modeling is done later in the process (schema-on-read).
- ETL might not be needed



## Ingest all Data

Extract and Load, NO Transform

## Store all data


In native format

## Do analysis

Using almost any tool

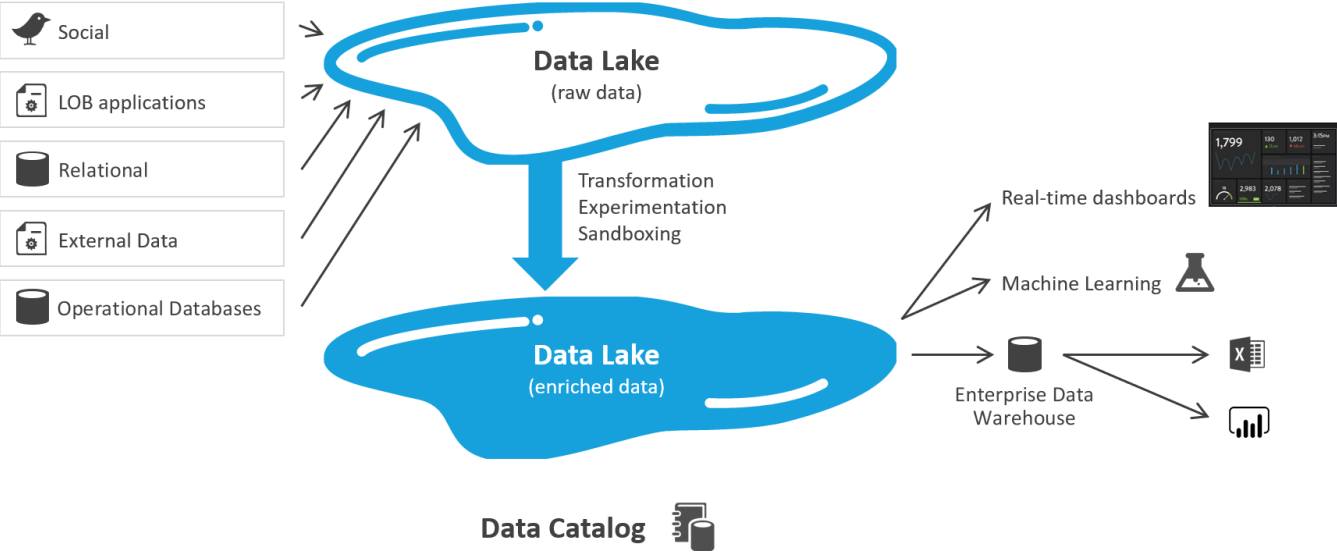

## Operationalize

Create schemas and pipelines

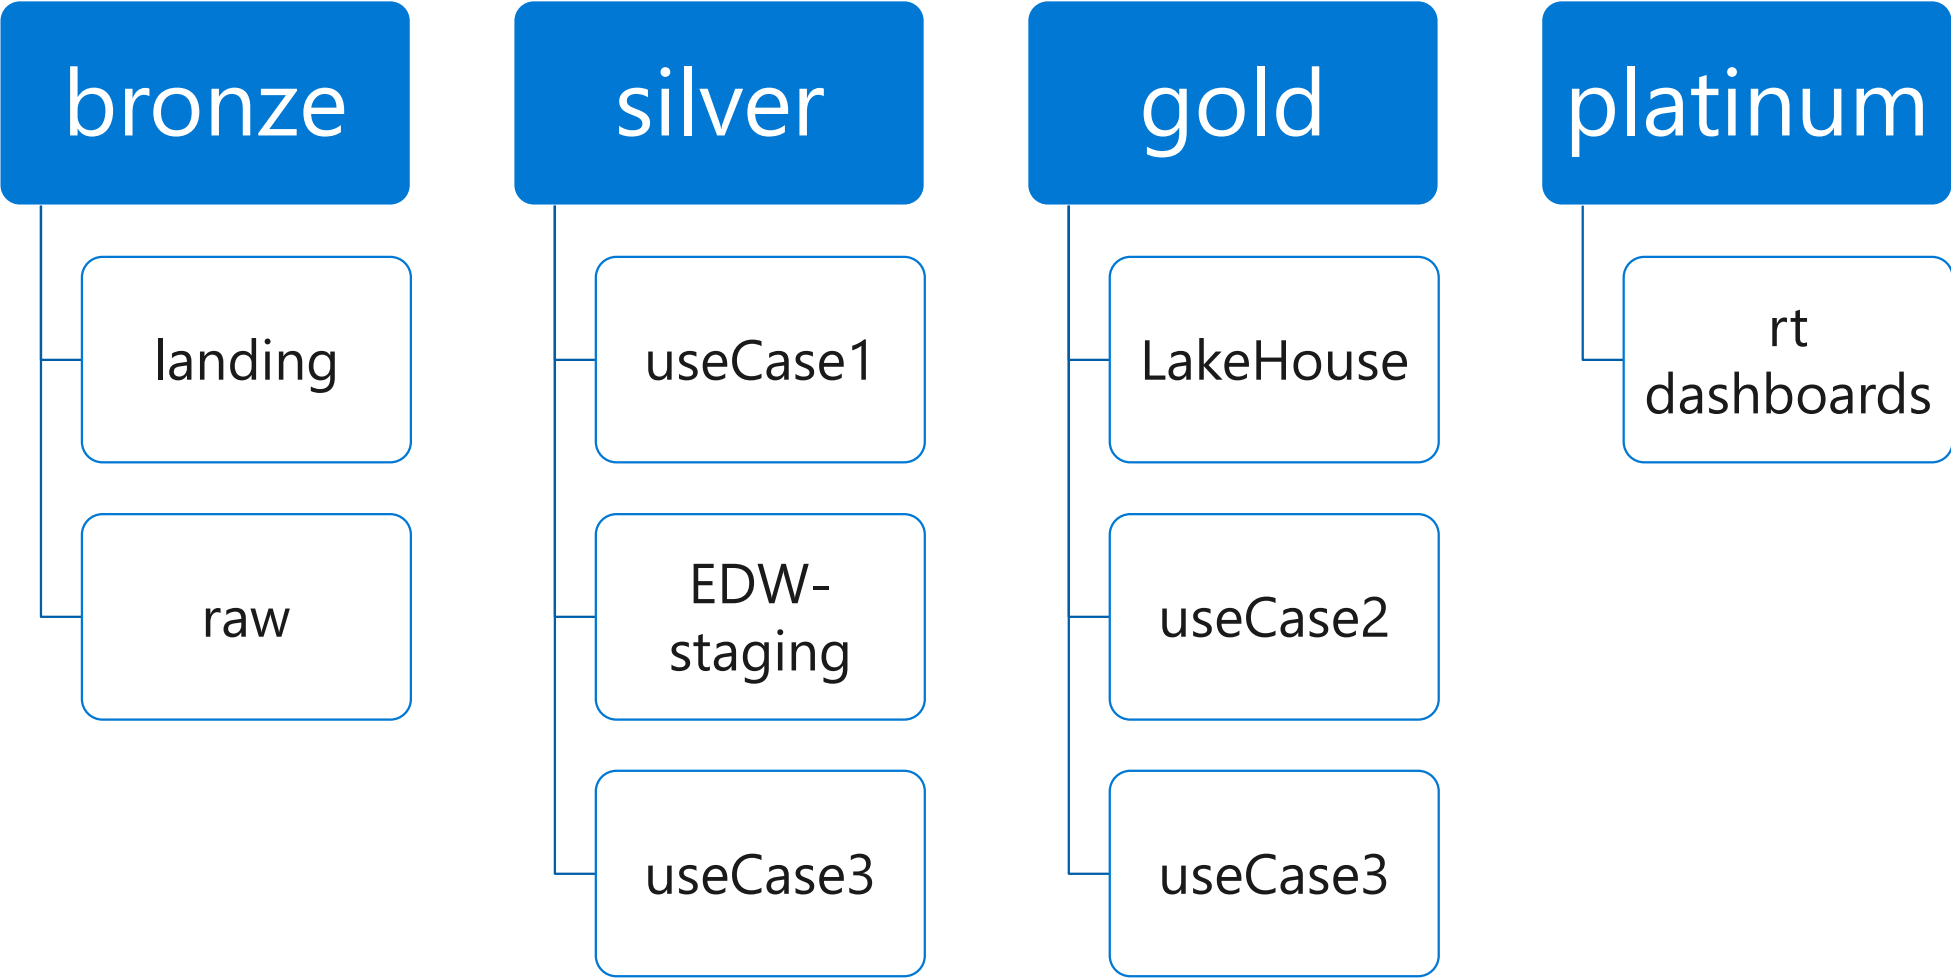




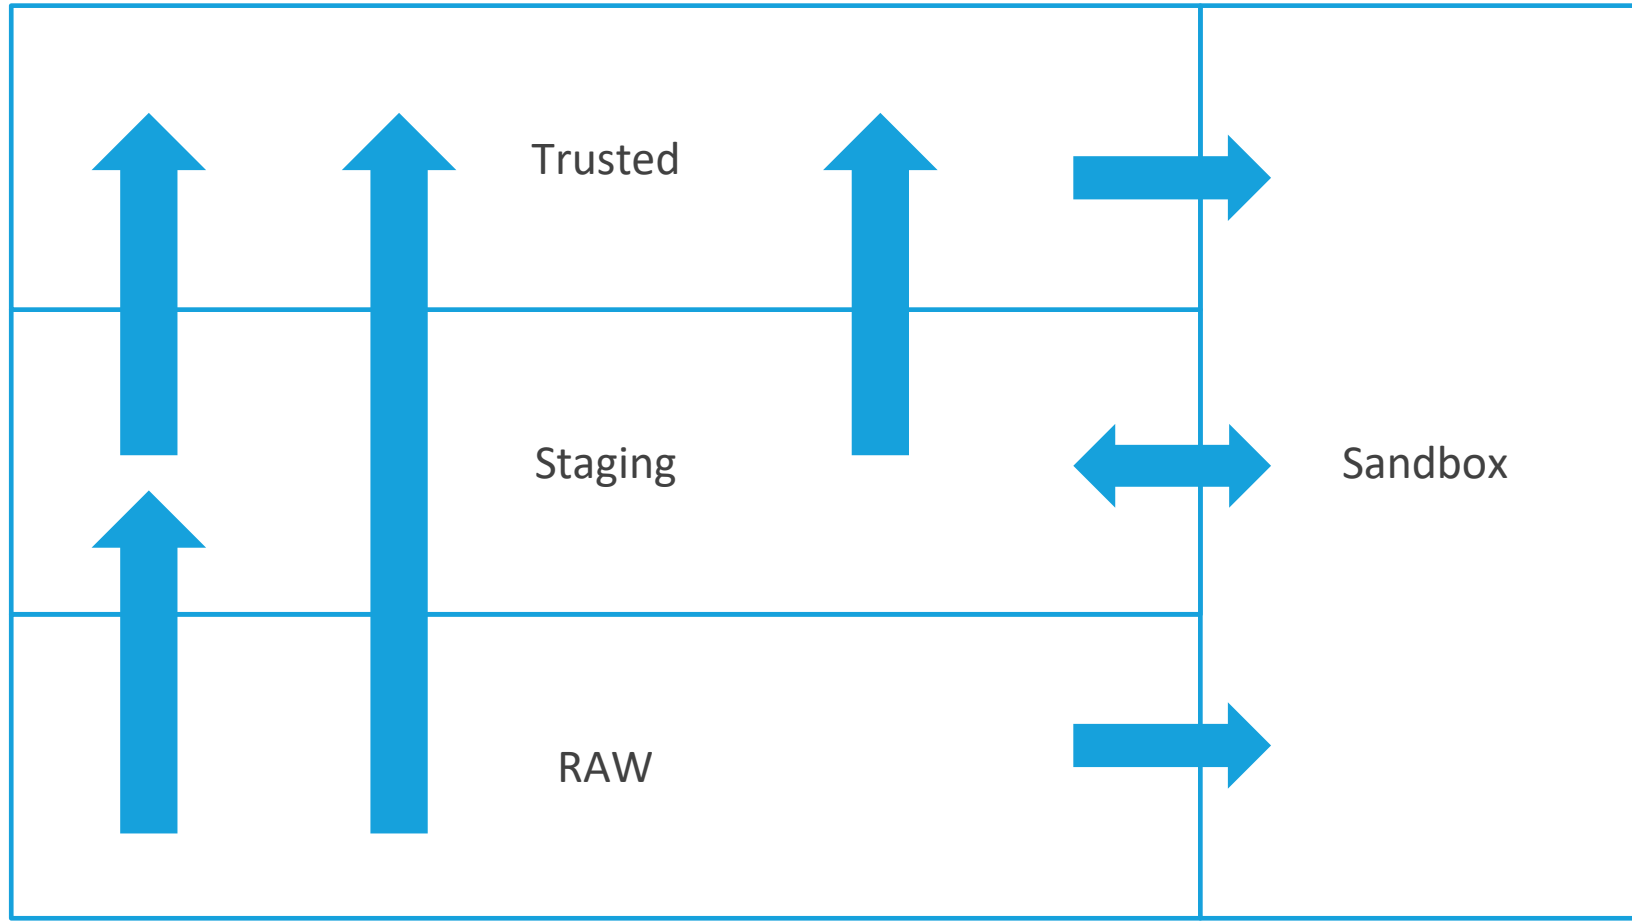
# Real-World Example – Customer 360



# Physical Structure



# Analytics workflow



- 1 Development / Data Ingestion
- 2 Production
- 3 Data Scientists / Business Analysts / Developers
- 4 Development
- 5 Production

# Data Lake vs. Data Warehouse

---

Data Lake	Data Warehouse
Complementary to the EDW	Can be sourced from the Data Lake
Load first, understand later	Understand first, load later
Schema-on-read	Schema-on-write
System of Insight	System of Record
Detailed Data	Refined Data
Supports varied data formats	Structured data
Adapts to changing requirements	Difficult to change structure
Optimized for Cost	Optimized for Performance

# Alternative Patterns to Traditional Data Warehouses

- Customer Data Platforms (CDP)
  - PLM Software
  - Demand Signal Repositories
  - Behavioral Data Platforms
  - traditional ERP/CRM software
- 
- Let's talk about the analytics patterns that work best



Search



Results Per Page: 10

Highlight

Open In ...

Delete

## Searches

## ▼ Current Search

[Save](#) | [Clear All](#)Source Type:  
Azure Data Lake Store

## ▼ Filters

## Tags

- ☐ earthquakes (4)
- ☐ datalaketest (3)
- ☐ Demo (2)
- ☐ datalaketest3 (2)

## Object Type

- ☐ Directory (3)
- ☐ File (3)
- ☐ Data Lake (2)

## Source Type

- ☐ SQL Server (4423)
- ☐ Filesystem (1551)
- ☐ SQL Server Analysis Services Tabular (396)
- ☐ SQL Server Analysis Services (140)

[see more](#)

## Daily Earthquake Records...



Directory represents logical data set. Contains one file per day.

Experts:  
**amitkul@microsoft.com**

earthquakes

Contained In Data Lake:  
**adlsadc.azuredatalakes...**

DATA LAKE DIRECTORY



Open In ... Explore Data Lake

## region\_names.asc

click tile to add a description...

Experts:  
**amitkul@microsoft.com**

earthquakes

Contained In Data Lake:  
**adlsadc.azuredatalakes...**

DATA LAKE FILE



Open In ... Explore Data Lake

## AzureDataCatalogErrorEve...

click tile to add a description...

Experts:  
**linm@test.com**

Demo datalaketest

## test1

click tile to add a description...

Experts:

datalaketest3

## &gt; Properties



Preview



Columns



Docs

Name:  
DailyEarthquakes

## Friendly Name:

Daily Earthquake Records from USGS

## Description:

Directory represents logical data set. Contains one file per day.

Experts:  
**amitkul@microsoft.com**

Add...

Tags:  
earthquakes

We have a  
winning idea, help  
us operationalize it

## Ingest all Data

Extract and Load, NO Transform

## Store all data

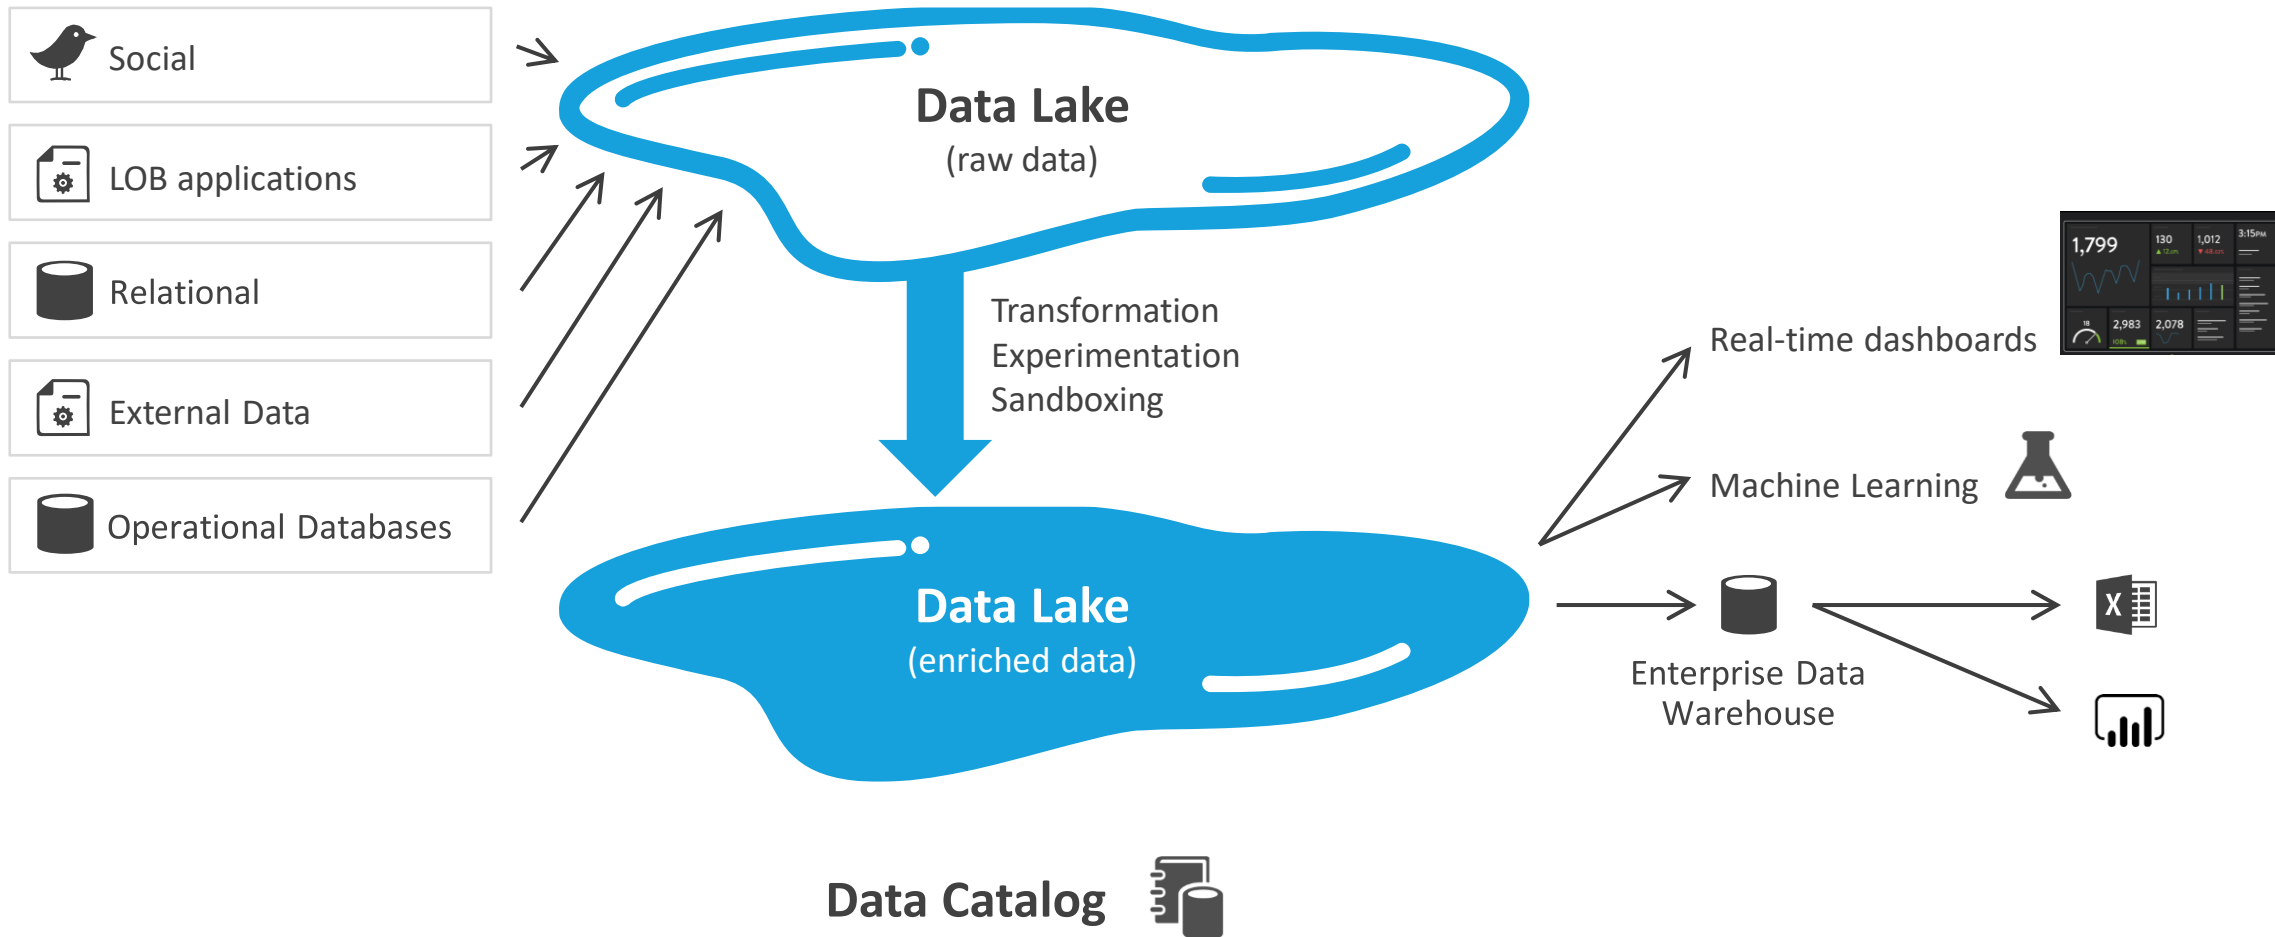
In native format

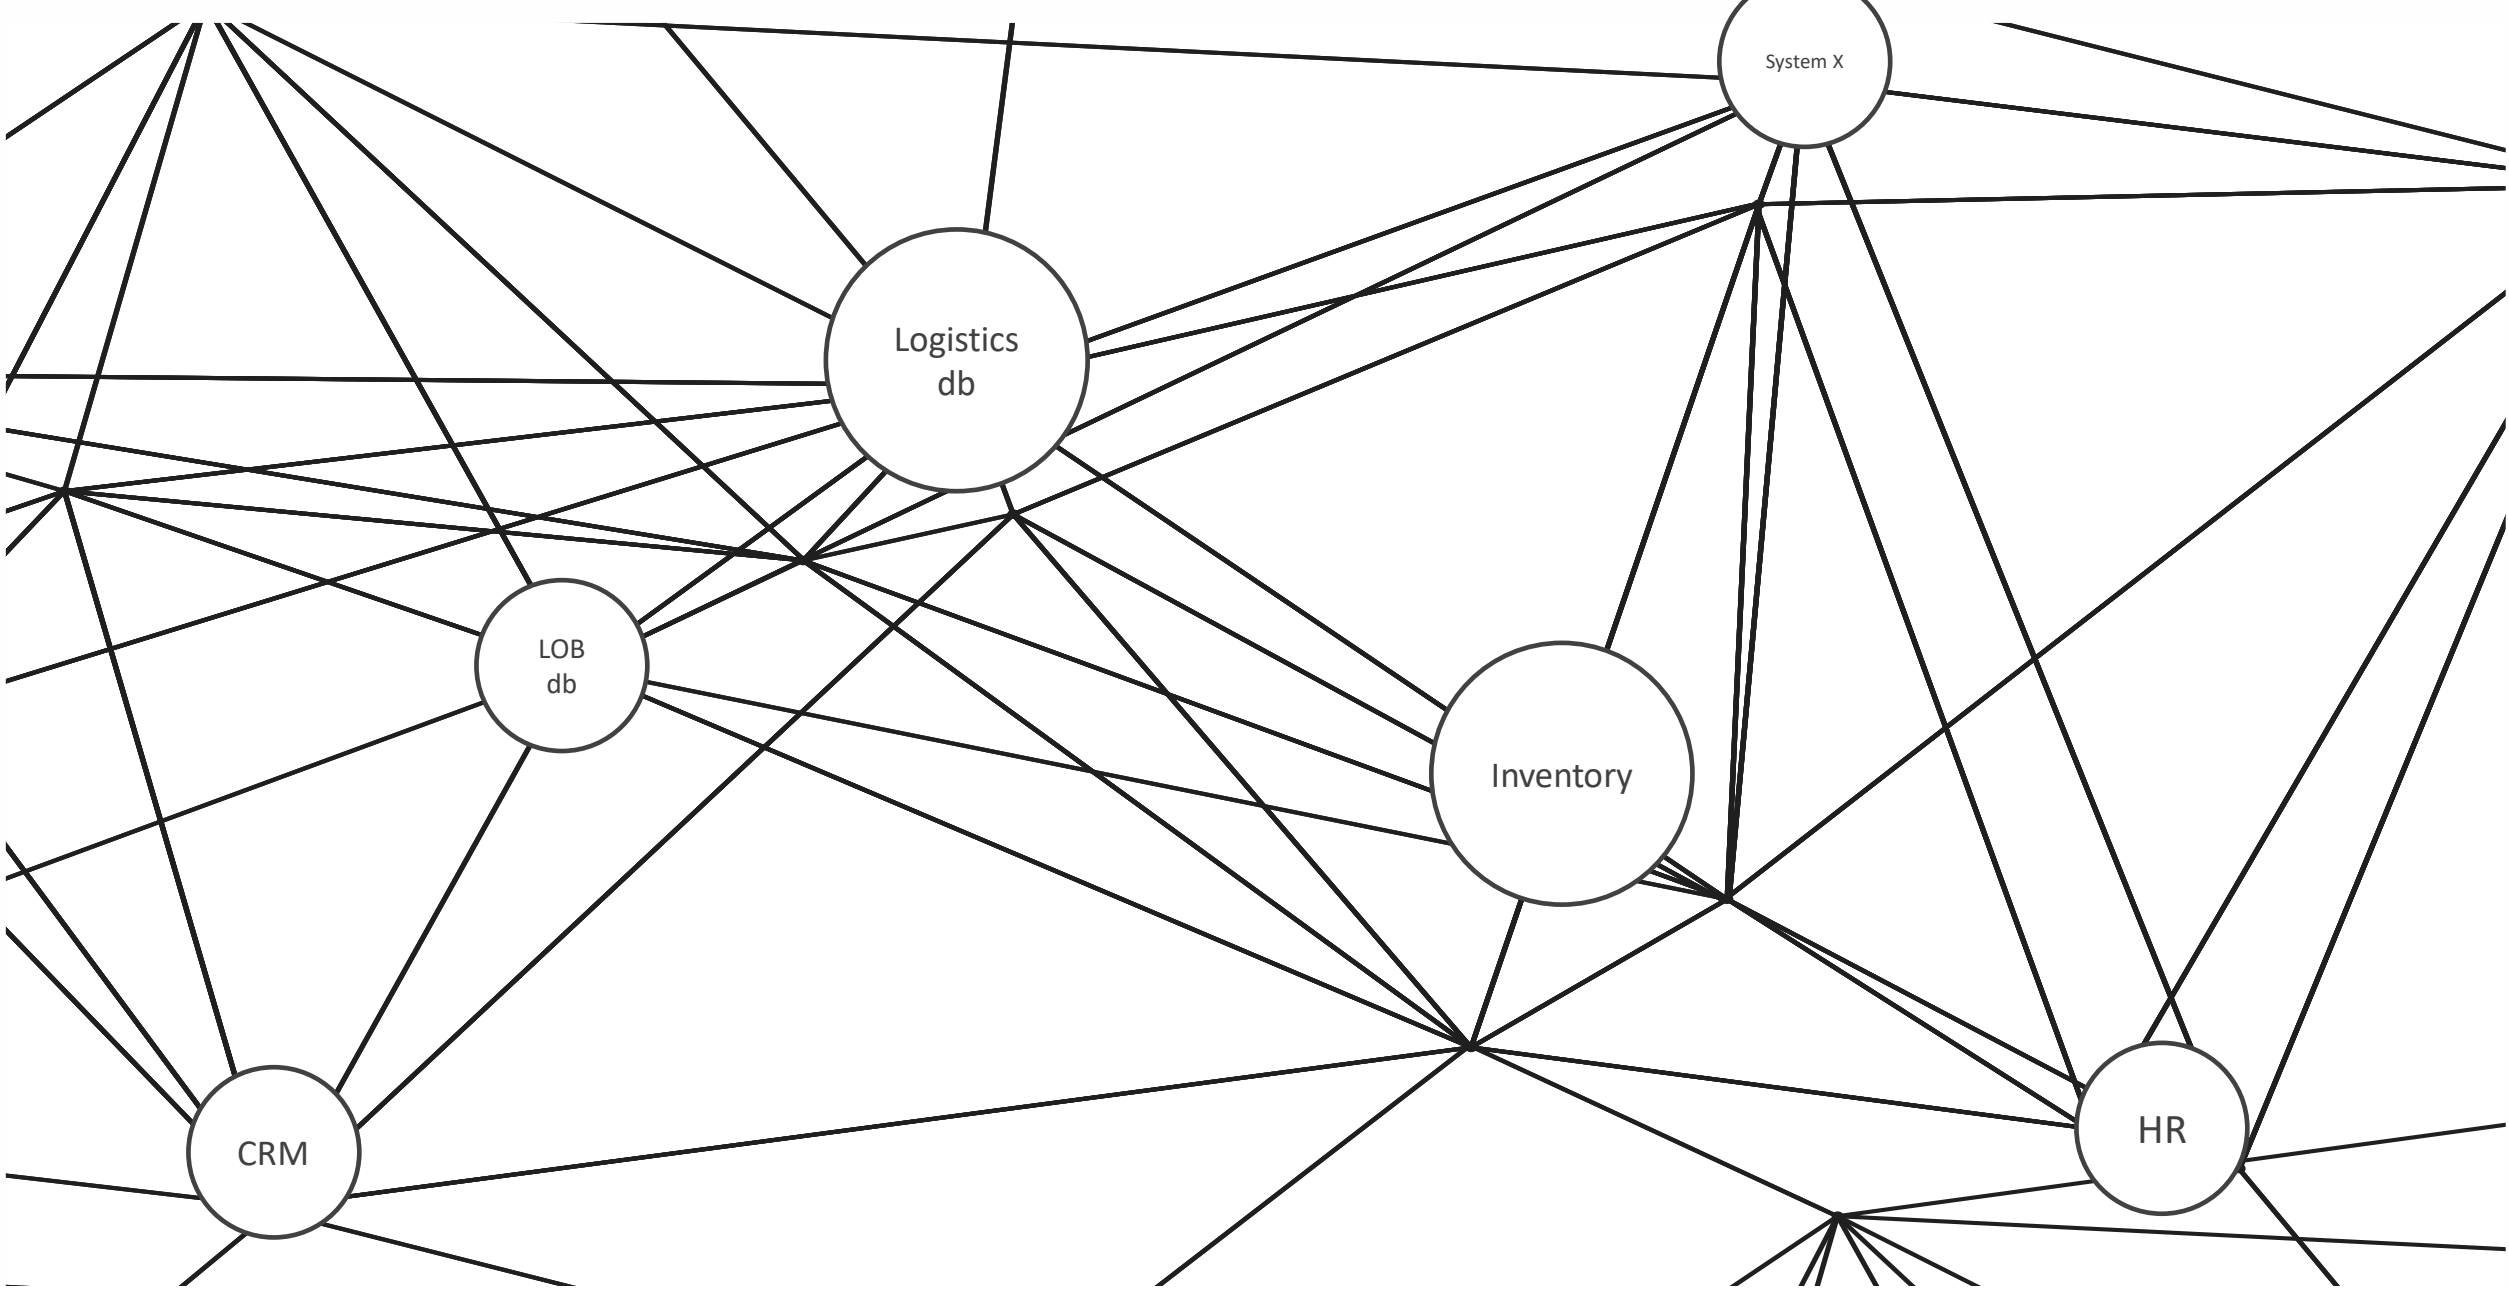
## Do analysis

Using almost any tool

## Operationalize

Create schemas and pipelines

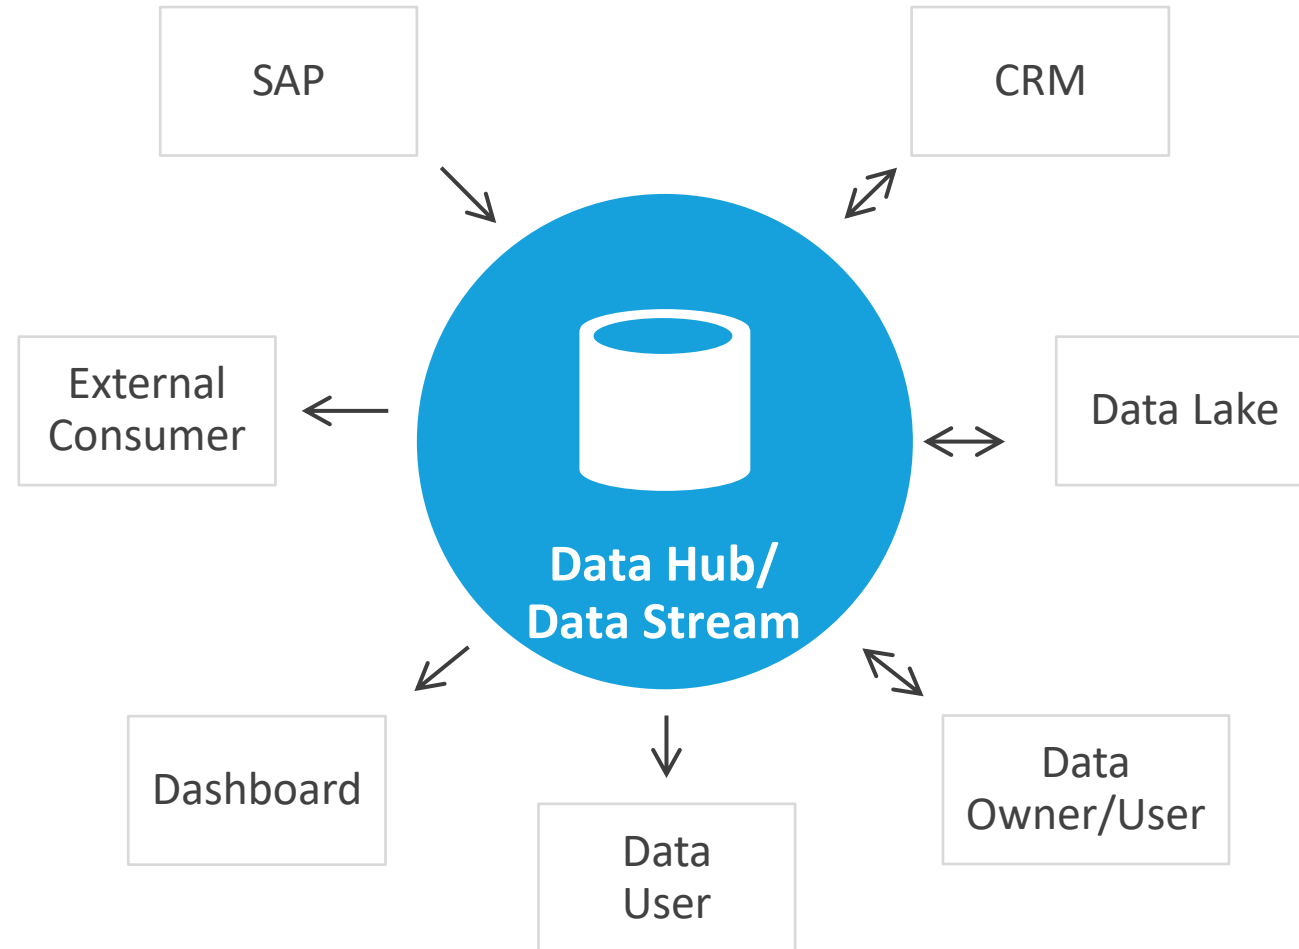




# Data Ingestion Factory


## (Kappa Architecture)

---





# Modern Data Warehouse



# Handling Streaming Data

